Jonathan Smith
ITAI 2372
A09: Ethical and Society Implications of AI Case Study

The Digital Lineup: A Comprehensive Report on Facial Recognition Technology in Law Enforcement

Introduction

Over the past decade, the landscape of criminal justice and law enforcement has been fundamentally altered by the rapid integration of artificial intelligence. Among the most powerful, controversial, and widely debated of these new tools is Facial Recognition Technology (FRT). What once existed strictly in the realm of science fiction is now an everyday reality in police departments across the globe. This report provides a detailed examination of facial recognition technology in the context of law enforcement. By exploring the mechanics of the AI system, its undeniable benefits, the profound ethical and societal implications it carries, and the inherent risks of its deployment, we can better understand how to govern its use. Furthermore, through practical recommendations and a detailed examination of the recent landmark *Williams v. City of Detroit* case, this report illustrates the profound human impact of algorithmic policing and the urgent need for comprehensive reform.

1. The AI System: Defining the Technology

Facial Recognition Technology is a highly sophisticated biometric artificial intelligence application capable of identifying or verifying a person’s identity from a digital image or a video frame. In the context of law enforcement, it serves as a digital detective designed to solve the persistent problem of slow, manual suspect identification. Traditionally, investigators relied on eyewitness accounts, physical fingerprints, or officers spending hundreds of tedious hours manually reviewing closed-circuit television (CCTV) footage to track a suspect's movements. FRT automates and exponentially accelerates this entire process.

At a high level, the technology operates through a four-step mechanical process:

Detection: A camera or software system scans an image or a video feed to detect the presence of a human face, isolating it from the background environment.

Analysis: The AI acts like a digital cartographer, mapping the specific facial features—often referred to as nodal points. It measures highly specific metrics, such as the exact distance between the eyes, the depth of the eye sockets, the shape of the cheekbones, and the contour of the lips.

Conversion: These physical measurements are translated into a complex mathematical formula, resulting in a unique digital signature known as a "faceprint," which is conceptually similar to a biometric fingerprint.

Matching: The system compares this newly generated faceprint against a vast database of known faces—typically compiled from police mugshots or state driver's license registries—to find a statistical match, thereby giving investigators a name to attach to an unknown face.

2. Benefits and Positive Impact

When operating under ideal conditions, the implementation of facial recognition technology offers extraordinary benefits to law enforcement agencies. The most immediate advantage is efficiency improvements. Investigations that previously took weeks or months of manual labor can now yield actionable leads in a matter of seconds. By rapidly narrowing down a suspect pool from thousands of individuals to a mere handful, the technology serves as a massive force multiplier for understaffed police departments.

Naturally, this efficiency translates directly into significant cost savings. By drastically reducing the human labor hours required to analyze video evidence or identify unknown suspects, police departments can allocate taxpayer funds and human resources far more effectively, shifting officers from desk work back into active community policing.

Furthermore, there are undeniable accuracy and innovation gains. Modern biometric systems can search through databases containing millions of faces much faster and, in some very specific operational environments, more accurately than human investigators. This is particularly useful when identifying prolific repeat offenders across different jurisdictions.

Ultimately, these technological advantages yield profound societal benefits. FRT has proven to be highly effective in time-sensitive, high-stakes situations where human lives are immediately on the line. The technology has been successfully deployed to locate missing children, identify victims of human trafficking who are unable to speak for themselves, and rapidly apprehend dangerous fugitives before they can cause further harm to the public.

3. Ethical Issues: The Moral Cost of Biometrics

Despite its operational benefits, the deployment of FRT has sparked intense ethical debates, primarily centering around bias and fairness. Machine learning algorithms are only as objective as the data on which they are trained. Historically, many facial recognition algorithms have been trained on datasets consisting predominantly of lighter-skinned males. For example, popular open-source training datasets, such as *Labeled Faces in the Wild*, have been found to consist of 83.5% white faces.

Consequently, these systems suffer from severe algorithmic bias. A landmark 2019 study by the National Institute of Standards and Technology (NIST) analyzed 189 facial recognition algorithms and found that many algorithms were 10 to 100 times more likely to misidentify a

Black or East Asian face compared to a white face. The algorithms were substantially less likely to correctly identify Black women than members of any other demographic group. In Detroit, the former Police Chief openly acknowledged that if the city's officers were to use facial recognition by itself, it would yield misidentifications 96% of the time. This leads to an unacceptably high risk of false identifications and wrongful arrests for people of color.

Beyond bias, the technology presents severe privacy violations. The use of FRT in public spaces via CCTV networks or police body cameras constitutes a form of mass, warrantless surveillance. It fundamentally infringes on the right to anonymity in public, allowing government entities to track where individuals travel, who they associate with, and what political or religious events they attend without any probable cause.

This directly ties into the issue of consent and data usage. FRT systems rely on massive databases built without explicit public consent. Citizens who simply apply for a driver's license at their local DMV do not consent to having their photo placed into a perpetual, virtual police lineup. Similarly, private AI companies have scraped billions of personal images from social media platforms to sell to law enforcement, completely bypassing user consent and commodifying personal data.

4. Societal Implications

The widespread integration of artificial intelligence in policing carries deep societal implications, beginning with the impact on jobs and the workforce. The nature of police work is shifting heavily toward digital forensics and technological analysis. While this creates a high demand for specialized AI operators and data analysts within law enforcement, it also requires extensive retraining for traditional detectives who may lack digital literacy, fundamentally changing the culture of policing.

More alarmingly, the technology has devastating effects on inequality and discrimination. Because surveillance technology is disproportionately deployed in heavily policed, lower-income, and minority neighborhoods, existing systemic inequalities are magnified. When the inherent algorithmic bias against darker skin tones is combined with the targeted deployment of cameras in minority communities, it creates a compounding discriminatory effect that effectively criminalizes certain demographics.

Ultimately, this unregulated deployment severely damages public trust and adoption. When communities feel they are being constantly watched and unfairly targeted by flawed machines, public trust in law enforcement evaporates. Widespread use of FRT creates a chilling effect on civil liberties, deterring innocent citizens from participating in lawful protests or public gatherings out of a very rational fear of being tracked, cataloged, and potentially falsely accused by an algorithm.

5. Risks and Challenges

The practical application of FRT is fraught with risks, starting with severe technical limitations. The technology frequently struggles in less-than-ideal conditions—a phenomenon often referred to as "garbage in, garbage out." Poor lighting, low-resolution security cameras, awkward camera angles, or subjects wearing masks, hats, or glasses can severely degrade the system's accuracy, turning a high-tech tool into a random guessing machine.

Furthermore, centralized biometric databases pose massive security risks. They are highly lucrative targets for state-sponsored hackers and cybercriminals. Unlike a compromised password or a stolen credit card number, you cannot "reset" or change your faceprint. If a law enforcement biometric database is breached, the victims' biometric identities are permanently compromised.

There are also severe misuse and abuse scenarios. The primary concern is "mission creep," where technology originally acquired to fight terrorism or severe violent crimes is eventually used to track minor infractions, such as jaywalking or petty theft. There is also the persistent risk of internal abuse, with documented cases of officers utilizing powerful surveillance systems to stalk ex-partners, harass journalists, or monitor political opponents without oversight.

6. Recommendations for Ethical Deployment

To harness the benefits of FRT while mitigating its harms, immediate and comprehensive reform is required.

Regarding policy and regulation, lawmakers must implement strict legal frameworks requiring law enforcement agencies to obtain a judge-approved warrant before utilizing facial recognition technology in an investigation. Furthermore, there should be a complete federal ban on "live" or real-time facial recognition on public cameras to prevent the realization of authoritarian-style mass surveillance infrastructure.

In terms of technical improvements, the government must mandate that any AI system used by law enforcement undergoes rigorous, independent auditing for demographic bias. Systems should be legally prohibited from deployment unless they can objectively demonstrate strict accuracy thresholds across all race and gender categories, effectively neutralizing the bias identified in the NIST studies.

Finally, agencies must adopt strict ethical guidelines. The most crucial of these is enforcing a strict "Human-in-the-Loop" policy. An AI-generated match should only ever be considered a preliminary investigative lead, not definitive evidence. It must never constitute probable cause for an arrest on its own. Additionally, strict data retention policies must be enforced, ensuring that the biometric data and images of innocent citizens are immediately purged from the system.

7. The Human Cost: Case Studies in Failure

To truly understand the dangers of unregulated AI in policing, one must look past the algorithms and examine the human toll. The case of *Williams v. City of Detroit* perfectly encapsulates the cascading failures of Facial Recognition Technology.

In 2018, police were attempting to identify a shoplifter who stole watches from a Shinola store. Utilizing only a grainy, low-quality still image from a security camera, the Detroit Police Department fed the blurry photo into an FRT system. The software ran the low-quality image against millions of state driver's license photos and flagged Robert Williams, a Black man, as a match.

In theory, the police believed the AI had instantly solved a cold case. In reality, it was a catastrophic failure highlighting bias and fairness. Williams was entirely innocent and looked nothing like the actual suspect other than having a similar physical build. Because the AI struggled to differentiate Black facial features—a known flaw—it delivered a false positive.

The officers suffered heavily from "automation bias," blindly trusting the computer's answer over basic human logic and investigative work. The detective exacerbated the situation through a lack of transparency, deliberately omitting from the arrest warrant application that the suspect identification was based solely on an unverified, low-quality FRT match. Furthermore, Williams' expired driver's license photo was used without his consent in this massive virtual lineup.

The societal implications were devastating. In January 2020, Williams was arrested on his front lawn in broad daylight, handcuffed in front of his wife and young children, and held in a detention center for 30 hours for a crime he did not commit. This trauma is not an isolated incident. Porcha Woodruff, another Black resident of Detroit, was falsely identified by FRT and arrested for a carjacking while she was eight months pregnant, getting her children ready for school. Cases like these destroy community trust and highlight the disproportionate harm inflicted on marginalized groups.

Fortunately, the Williams case sparked real-world change. In June 2024, a landmark civil rights lawsuit brought by Williams was settled, resulting in the nation's strictest police policy for FRT. The new Detroit Police Department policy enforces the exact recommendations experts have pleaded for: it mandates a strict "Human-in-the-Loop" framework where an FRT match can never be the sole basis for an arrest, and it explicitly bans police from running FRT searches on low-quality or blurry images. Robert Williams' nightmare forced the creation of a blueprint for responsible AI governance, proving that technology must always be subordinate to human rights.

Works Cited

American Civil Liberties Union. "After Third Wrongful Arrest, ACLU Slams Detroit Police Department for Continuing to Use Faulty Facial Recognition Technology." ACLU, 6 Aug. 2023, www.aclu.org/press-releases/after-third-wrongful-arrest-aclu-slams-detroit-police-department-for-continuing-to-use-faulty-facial-recognition-technology. Accessed 1 Apr. 2026.

American Civil Liberties Union. "Williams v. City of Detroit." ACLU, 28 Jun. 2024, www.aclu.org/cases/williams-v-city-of-detroit-face-recognition-false-arrest. Accessed 1 Apr. 2026.

Cwiek, Sarah. "“It didn't make sense at all”: Wrongful facial recognition arrest in Detroit leads to landmark settlement." Michigan Public, 28 Jun. 2024, www.michiganpublic.org/criminal-justice-legal-system/2024-06-28/it-didnt-make-sense-at-all-wrongful-facial-recognition-arrest-leads-to-landmark-settlement. Accessed 1 Apr. 2026.

Harvard Journal of Law & Technology. "Why Racial Bias is Prevalent in Facial Recognition Technology." JOLT Digest, 3 Nov. 2020, jolt.law.harvard.edu/digest/why-racial-bias-is-prevalent-in-facial-recognition-technology. Accessed 1 Apr. 2026.

Swarns, Christina. "When Artificial Intelligence Gets It Wrong." Innocence Project, 19 Sep. 2023, innocenceproject.org/news/when-artificial-intelligence-gets-it-wrong/. Accessed 1 Apr. 2026.

United States, National Institute of Standards and Technology. Face Recognition Vendor Test (FRVT) Part 8: Summarizing Demographic Differentials. NIST Interagency Report 8429, 2022, nvlpubs.nist.gov/nistpubs/ir/2022/NIST.IR.8429.ipd.pdf. Accessed 1 Apr. 2026.

Williams, Robert. "I Was Wrongfully Arrested Because of Facial Recognition Technology. It Shouldn't Happen to Anyone Else." ACLU of Michigan / TIME Magazine, 11 Jul. 2024, www.aclumich.org/news/i-was-wrongfully-arrested-because-facial-recognition-technology-it-shouldnt-happen-anyone-else/. Accessed 1 Apr. 2026.